

One point per defect, at its strongest real-world measure

real — do defective and clean code differ on it?

injected — does changing only the defect move the measure?

CO-OCCURRENCE
the measure provably cannot respond,
yet real defective code separates anyway

deep nesting
(cyclomatic)

equal on both sources

complex conditional
(halstead difficulty)

long method
(halstead difficulty)

inefficient copy (cognitive)

duplicate code (sloc)

both numbers are the same size artefact

magic number
(maintainability)

inefficient loop (cyclomatic)

broad except (maintainability)

long parameter list (halstead difficulty)

dead code (api calls)

mutable default (comment density)

perf try in loop (api calls)

BLIND
neither source

DETECTION
responds on both sources

verdict

- detects
- co-occurs
- blind

Detected defects sit ABOVE the diagonal, and that is expected: real defective code differs from clean code in more ways than the defect alone, so the real value is inflated. The diagonal is a reference, not a target.