

Scoring raw output partly measures verbosity:
one model's rate halves once its own test cases are excluded

